

[LIVE SESSION] How I Landed my Dream Data Science Job? | 13th May

Register Now (https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

 LOGIN / REGISTER ([HTTPS://ID.ANALYTICSVIDHYA.COM/](https://id.analyticsvidhya.com/))

/04/FEAT



Machine Learning
— 50 Most C

Do

(<https://www.analyticsvidhya.com/back-channel/>)



(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

BEGINNER ([HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/CATEGORY/BEGINNER/](https://www.analyticsvidhya.com/blog/category/beginner/))

DATA ENGINEERING ([HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/CATEGORY/DATA-ENGINEERING/](https://www.analyticsvidhya.com/blog/category/data-engineering/))

PYTHON ([HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/CATEGORY/PYTHON-2/](https://www.analyticsvidhya.com/blog/category/python-2/))

REGRESSION ([HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/CATEGORY/REGRESSION/](https://www.analyticsvidhya.com/blog/category/regression/))

STRUCTURED DATA ([HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/CATEGORY/STRUCTURED-DATA/](https://www.analyticsvidhya.com/blog/category/structured-data/))

TECHNIQUE ([HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/CATEGORY/TECHNIQUE/](https://www.analyticsvidhya.com/blog/category/technique/))

Feature Scaling for Machine Learning: Understanding the Difference Between Normalization vs. Standardization

ANIRUDDHA BHANDARI ([HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/AUTHOR/ANIRUDDHA/](https://www.analyticsvidhya.com/blog/author/aniruddha/)), APRIL 3, 2020 [LOGIN TO BOOKMARK THIS ARTICLE](#)

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our [Privacy Policy](https://www.analyticsvidhya.com/privacy-policy/)

(<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Introduction to Feature Scaling

Accept

I was recently working with a dataset that had multiple features spanning varying degrees of magnitude, range, and units. This is a significant obstacle as a few machine learning algorithms are highly sensitive to these features.

I'm sure most of you must have faced this issue in your dataset where one feature is entirely in kilograms while the other is in grams. This is a problem when features vary so vastly in terms of what they represent.

This is where I turned to the concept of data preprocessing stage but I've realized that it's at the detriment of their machine learning models.

Analytics Vidhya

Webinar

How I Landed my Dream Data Science Job?

13th May 2021

7.30 PM - 8.30 PM

REGISTER NOW

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

(<https://cdn.analyticsvidhya.com/wp-content/uploads/2020/04/Feature-image-Normalization-vs-Standardization.png>)

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya you agree to our [Privacy Policy](https://www.analyticsvidhya.com/privacy-policy/) and [Terms of Use](https://www.analyticsvidhya.com/terms/). What could be the reason behind this quirk?

Also, what's the difference between normalization and standardization? These are two of the most commonly used feature scaling techniques in machine learning but a level of ambiguity exists in their understanding. When

should you use which technique?

I will answer these questions and more in this article on feature scaling. We will also implement feature scaling in Python to give you a practice understanding of how to do it.

Note: I assume that you are familiar with Python and I recommend going through the below courses:

- [Python for Data Science \(https://courses.analyticsvidhya.com/courses/python-for-data-science?utm_source=blog&utm_medium=feature-scaling\)](https://courses.analyticsvidhya.com/courses/python-for-data-science?utm_source=blog&utm_medium=feature-scaling)
- [All free Machine Learning Courses by Analytics Vidhya \(https://courses.analyticsvidhya.com/collections?category=free\)](https://courses.analyticsvidhya.com/collections?category=free)
- [Applied Machine Learning \(https://courses.analyticsvidhya.com/courses/applied-machine-learning-beginner-to-professional?utm_source=blog&utm_medium=feature-scaling\)](https://courses.analyticsvidhya.com/courses/applied-machine-learning-beginner-to-professional?utm_source=blog&utm_medium=feature-scaling)

Table of Contents

1. Why Should we Use Feature Scaling?
2. What is Normalization?
3. What is Standardization?
4. The Big Question – Normalize or Standardize?
5. Implementing Feature Scaling in Python
 - Normalization using Sklearn
 - Standardization using Sklearn
6. Applying Feature Scaling to Machine Learning Algorithms
 - K-Nearest Neighbours (KNN)
 - Support Vector Regressor
 - Decision Tree

Why Should we Use Feature Scaling?

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our [Privacy Policy \(https://www.analyticsvidhya.com/privacy-policy/\)](https://www.analyticsvidhya.com/privacy-policy/) and [Terms of Use \(https://www.analyticsvidhya.com/terms/\)](https://www.analyticsvidhya.com/terms/). The first question we need to address – why do we need to scale the variables in our dataset? Some machine learning algorithms are sensitive to feature scaling while others are virtually invariant to it. Let me explain that in more detail.

Accept

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

Gradient Descent Based Algorithms

Machine learning algorithms like [linear regression](https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization) [network-from-scratch-in-python-and-r/?utm_source=normalization-standardization](https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization)), [logistic regression](https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization) [network-from-scratch-in-python-and-r/?utm_source=normalization-standardization](https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization)), [neural network](https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization) ([network-from-scratch-in-python-and-r/?utm_source=normalization-standardization](https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization)), etc. that use gradient descent to find the minima of the cost function. The features should be scaled. Take a look at the formula for gradient descent:

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (y_i - \hat{y}_i) x_i$$

(<https://cdn.analyticsvidhya.com/wp-content/uploads/2019/08/gradient-descent-formula.png>)

The presence of feature value X in the formula will cause features with different ranges of features will cause different step sizes for the gradient descent to converge smoothly towards the minima and that the steps for gradient descent are updated at the same rate for all the features, we scale the data before feeding it to the model.

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

Having features on a similar scale can help the gradient descent converge more quickly towards the minima.

Distance-Based Algorithms

Distance algorithms like [KNN](https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization) (https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization), [K-means](https://www.analyticsvidhya.com/blog/2019/08/comprehensive-guide-k-means-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization) (https://www.analyticsvidhya.com/blog/2019/08/comprehensive-guide-k-means-clustering/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization),

[and SVM](https://www.analyticsvidhya.com/privacy-policy/) (<https://www.analyticsvidhya.com/privacy-policy/>) and [Terms of Use](https://www.analyticsvidhya.com/terms/) (<https://www.analyticsvidhya.com/terms/>).
code/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization) are

Accept

most affected by the range of features. This is because behind the scenes they are using distances between data points to determine their similarity.

For example, let's say we have data containing high their future incomes (in thousands Rupees):

Student	
0	
1	
2	
3	
4	

(https://cdn.analyticsvidhya.com/wp-content/uploads/2020/03/knn_ex_scaled.png)

Since both the features have different scales, there is a higher magnitude. This will impact the performance. We want our algorithm to be biased towards one feature.

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

Therefore, we scale our data before employing a distance based algorithm so that all the features contribute equally to the result.

	Student	CGPA	Salary '000
0	1	-1.184341	1.520013
1	2	-1.184341	-1.100699
2	3	0.416120	-1.100699
3	4	1.216350	0.209657
4	5	0.736212	0.471728

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (https://cdn.analyticsvidhya.com/wp-content/uploads/2020/03/knn_ex_scaled.png) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

The effect of scaling is conspicuous when we compare the Euclidean distance between data points for students

A and B, and between B and C, before and after scaling as shown below:

- Distance AB before scaling $\Rightarrow \sqrt{(40 - 60)^2 + (3 - 3)^2} = 20$
- Distance BC before scaling $\Rightarrow \sqrt{(40 - 40)^2 + (3 - 3)^2} = 0$
- Distance AB after scaling $\Rightarrow \sqrt{(1.1 + 1.5)^2 + (1.1 - 1.1)^2} = 2.6$
- Distance BC after scaling $\Rightarrow \sqrt{(1.1 - 1.1)^2 + (1.1 - 1.1)^2} = 0$

Scaling has brought both the features into the picture where they were before we applied scaling.

Tree-Based Algorithms

Tree-based algorithms (https://www.analyticsvidhya.com/tutorial-scratch-in-python/?utm_source=blog&utm_medium=popup), on the other hand, are fairly insensitive to the scale of the features. A decision tree is only splitting a node based on a single feature and the homogeneity of the node. This split on a feature

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

So, there is virtually no effect of the remaining features on the split. This is what makes them invariant to the scale of the features!

What is Normalization?

Normalization is a scaling technique in which values are shifted and rescaled so that they end up ranging between 0 and 1. It is also known as Min-Max scaling.

Here's the formula for normalization:

$$X' = \frac{X - X_{min}}{X_{max} - X_{min}}$$

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (https://cdn.analyticsvidhya.com/wp-content/uploads/2020/03/Norm_eq.gif) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Here, X_{max} and X_{min} are the maximum and the minimum values of the feature respectively.

- When the value of X is the minimum value in the column, the numerator will be 0, and hence X' is 0
- On the other hand, when the value of X is the maximum value in the column, the numerator is equal to the denominator and thus the value of X' is 1
- If the value of X is between the minimum and

What is Standardization?

Standardization is another scaling technique where standard deviation. This means that the mean of the has a unit standard deviation.

Here's the formula for standardization:

$$X'$$

(<https://cdn.analyticsvidhya.com/wp-content/uploa>

μ is the mean of the feature values and σ is the standard deviation. The values are not restricted to a particular range.

Now, the big question in your mind must be when should we use normalization and when should we use standardization? Let's find out!

The Big Question – Normalize or Standardize?

Normalization vs. standardization is an eternal question among machine learning newcomers. Let me elaborate on the answer in this section.

- Normalization is good to use when you know that the distribution of your data does not follow a Gaussian distribution. This can be useful in algorithms that do not assume any distribution of the data like K-Nearest Neighbors and Neural Networks.

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy.

However, this does not have to be necessarily true. Also, unlike normalization, standardization does not have a bounding range. So, even if you have outliers in your data, they will not be affected by

standardization.

Accept

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

However, at the end of the day, the choice of using normalization or standardization will depend on your problem and the machine learning algorithm you are using. There is no hard and fast rule to tell you when to normalize or standardize your data. You can always start by fitting your model to raw, normalized and standardized data and compare the performance for best results.

It is a good practice to fit the scaler on the training data to avoid any data leakage during the model testing process required.

Implementing Feature Scaling in Python

Now comes the fun part – putting what we have learned about machine learning algorithms on the [Big Mart dataset](https://www.analyticsvidhya.com/blog/2020/04/problem-big-mart-sales-iii/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization) (https://www.analyticsvidhya.com/blog/2020/04/problem-big-mart-sales-iii/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization). I've taken the [DataHack](https://www.analyticsvidhya.com/blog/2020/04/datahack/) (https://www.analyticsvidhya.com/blog/2020/04/datahack/) on the hackathon leaderboard (so that's worth checking out!)

I will skip the preprocessing steps since they are out of scope and explained in this [article](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/) (https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/).

Those steps will enable you to reach the top 20 percentile on the hackathon leaderboard (so that's worth checking out!)

So, let's first split our data into training and testing sets:

```
1 # splitting training and testing data
2 from sklearn.model_selection import train_test_split
3
4 X = df
5 y = target
6
7 X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=27)
```

[view raw](https://gist.github.com/aniruddha27/02432c8623b5e202f79c3964bd559867) (https://gist.github.com/aniruddha27/02432c8623b5e202f79c3964bd559867)

[/raw/a423c2def59d31259b0b1e358c0700112a4b42e1/NormalizationVsStandarization_1.py](https://gist.github.com/aniruddha27/02432c8623b5e202f79c3964bd559867#file-NormalizationVsStandarization_1.py)

[NormalizationVsStandarization_1.py](https://gist.github.com/aniruddha27/02432c8623b5e202f79c3964bd559867#file-NormalizationVsStandarization_1.py) (https://gist.github.com/aniruddha27/02432c8623b5e202f79c3964bd559867#file-NormalizationVsStandarization_1.py)

We use cookies on Analytics Vidhya website to enhance navigation, analyze site usage, and assist in our marketing efforts. (https://www.analyticsvidhya.com/privacy-policy) and Terms of Use (https://www.analyticsvidhya.com/terms/).

experience on the site. By using Analytics Vidhya, you agree to our [Privacy Policy](https://www.analyticsvidhya.com/privacy-policy)

Before moving to the feature scaling part, let's glance at the details about our data using the `describe()` method:

Accept

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Outlet_Size	Tier 2	Tier 3	Supermarket Type1	Supermarket Type2
count	6818.000000	6818.000000	6818.000000	6818.000000	6818.000000	6818.000000	6818.000000	6818.000000	6818.000000	6818.000000
mean	12.835420	0.355676	-2.940445	140.486413	15.151884	0.820202	0.226010	0.205571	0.651071	0.111616
std	4.233450	0.478753	0.791551	62.067053	1.151884	0.220202	0.226010	0.205571	0.651071	0.111616
min	4.555000	0.000000	-5.633875	31.290000	11.151884	0.000000	0.000000	0.000000	0.000000	0.000000
25%	9.300000	0.000000	-3.467944	93.385700	12.151884	0.000000	0.000000	0.000000	0.000000	0.000000
50%	12.857645	0.000000	-2.862535	142.179900	13.151884	0.000000	0.000000	0.000000	0.000000	0.000000
75%	16.000000	1.000000	-2.331264	184.495000	14.151884	0.000000	0.000000	0.000000	0.000000	0.000000
max	21.350000	1.000000	-1.113550	266.888400	15.151884	0.820202	0.226010	0.205571	0.651071	0.111616

(https://cdn.analyticsvidhya.com/wp-content/uploads/2019/04/supermarket_data.csv)

We can see that there is a huge difference in the range of values for Item_Visibility, Item_Weight, Item_MRP, and Outlet_Establishment_Year. This is why we need feature scaling!

Note: You will notice negative values in the Item_Visibility feature. This is because of the skewness in the feature.

Normalization using sklearn

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

To normalize your data, you need to import the *MinMaxScaler* from the *sklearn*

(https://courses.analyticsvidhya.com/courses/get-started-with-scikit-learn-sklearn?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization) library and apply it to our dataset. So, let's do that!

```

1 # data normalization with sklearn
2 from sklearn.preprocessing import MinMaxScaler
3
4 # fit scaler on training data
5 norm = MinMaxScaler().fit(X_train)
6
7 # transform training data
8 X_train_norm = norm.transform(X_train)
9
10 # transform testing data
11 X_test_norm = norm.transform(X_test)

```

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Accept

view raw (<https://gist.github.com/aniruddha27/a41a35725ec02006b05156e5483cb184>)

/raw/5a15a3fbae1a6967171185127c458674cd021f22/NormalizationVsStandarization_2.py)

NormalizationVsStandarization_2.py (https://gist.github.com/aniruddha27/a41a35725ec02006bb3156e5483cb184#file-normalizationvsstandarization_2-py) hosted with ❤ by GitHub (<https://github.com>)


Let's see how normalization has affected our dataset

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_MRP	Outlet
count	6818.000000	6818.000000	6818.000000	6818.000000	
mean	0.493029	0.355676	0.595849	0.463485	
std	0.252066	0.478753	0.175109	0.263444	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.282525	0.000000	0.479154	0.263566	
50%	0.494352	0.000000	0.613084	0.470673	
75%	0.681453	1.000000	0.730614	0.650280	
max	1.000000	1.000000	1.000000	1.000000	

(<https://cdn.analyticsvidhya.com/wp-content/uploads/2020/04/feature-scaling-machine-learning-models.pdf>)

All the features now have a minimum value of 0 and

Try out the above code in the live coding window below



Analytics Vidhya

Webinar

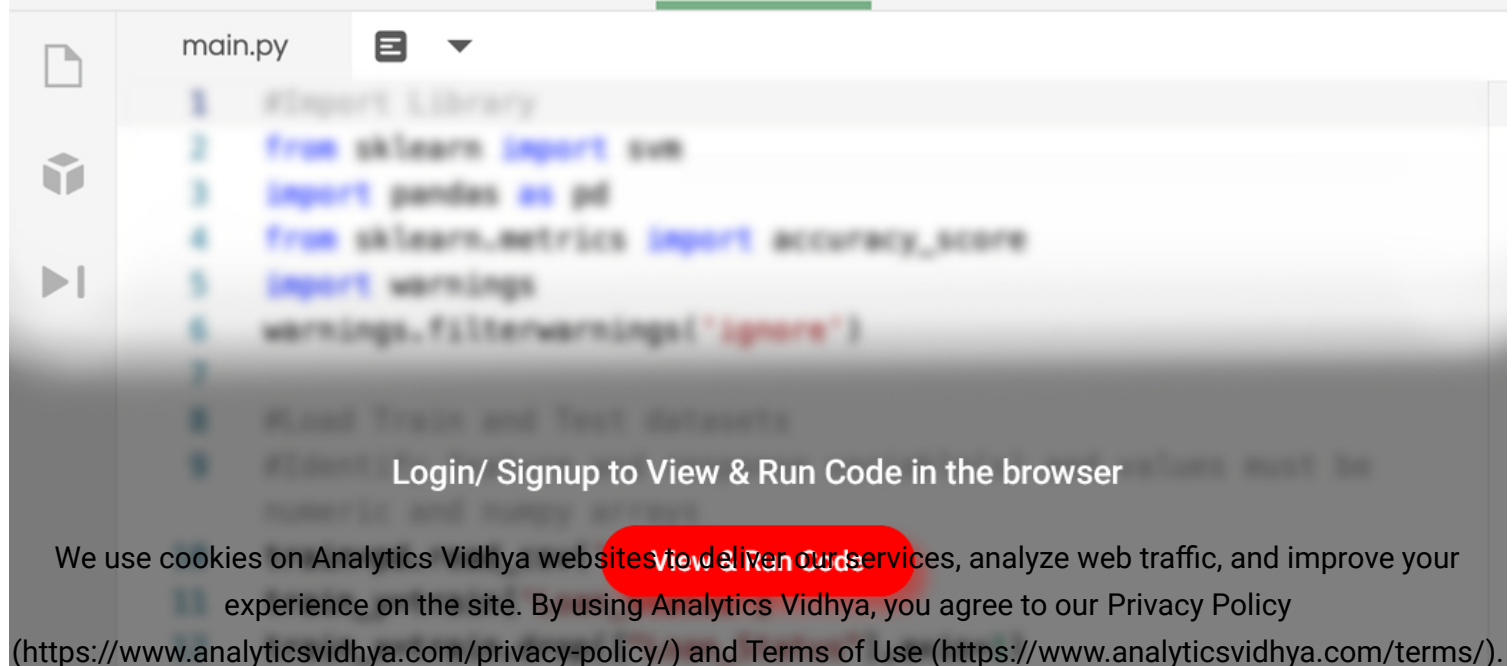
How I Landed my Dream Data Science Job?

13th May 2021

7.30 PM - 8.30 PM

REGISTER NOW

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)



main.py

```

1 # Import Library
2 from sklearn import svm
3 import pandas as pd
4 from sklearn.metrics import accuracy_score
5 import warnings
6 warnings.filterwarnings("ignore")
7
8 # Load Train and Test datasets
9
10 # Train the model

```

Login/ Signup to View & Run Code in the browser

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Model Code

Accept

(<https://id.analyticsvidhya.com/auth/login/?next=https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-models.pdf>)

[scaling-machine-learning-normalization-standardization/?utm_source=coding-window-blog&source=coding-window-blog\)](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-...)

Next, let's try to standardize our data.

Standardization using sklearn

To standardize your data, you need to import the `StandardScaler` dataset. Here's how you can do it:

```
1 # data standardization with sklearn
2 from sklearn.preprocessing import StandardScaler
3
4 # copy of datasets
5 X_train_stand = X_train.copy()
6 X_test_stand = X_test.copy()
7
8 # numerical features
9 num_cols = ['Item_Weight', 'Item_Visibility', 'Item_MRP', 'Outlet_Establishment_Year']
10
11 # apply standardization on numerical features
12 for i in num_cols:
13
14     # fit on training data column
15     scale = StandardScaler().fit(X_train_stand[[i]])
16
17     # transform the training data column
18     X_train_stand[i] = scale.transform(X_train_stand[[i]])
19
20     # transform the testing data column
21     X_test_stand[i] = scale.transform(X_test_stand[[i]])
```

view raw (https://gist.github.com/aniruddha27/965ff8b01e19de1cffdb5cbe703d5495/raw/948e4d05a4e65291ff54f6986f552ea398af89de/NormalizationVsStandardization_3.py)

We use cookies on Analytics Vidhya websites to deliver our services, enhance user navigation, analyze site usage, and assist in our marketing efforts. (https://www.analyticsvidhya.com/privacy-policy/) and Terms of Use (https://www.analyticsvidhya.com/terms/). You would have noticed that I only applied standardization to my numerical columns and not the other One-Hot Encoded (<https://www.analyticsvidhya.com/blog/2020/03/one-hot-encoding-vs-label-encoding-using-scikit-learn>)

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

[learn/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/)) features. Standardizing the One-Hot encoded features would mean assigning a distribution to categorical features. You don't want to do that!

But why did I not do the same while normalizing the range between 0 to 1. So, normalization would not a

Right, let's have a look at how standardization has tr

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_MRP	Out
count	6.818000e+03	6818.000000	6.818000e+03	6.818000e+03	
mean	1.704754e-16	0.355676	2.342737e-16	1.233002e-16	
std	1.000073e+00	0.478753	1.000073e+00	1.000073e+00	
min	-1.956094e+00	0.000000	-3.402972e+00	-1.759459e+00	
25%	-8.351767e-01	0.000000	-6.664603e-01	-7.589239e-01	
50%	5.250371e-03	0.000000	9.843446e-02	2.728679e-02	
75%	7.475728e-01	1.000000	7.696596e-01	7.091011e-01	
max	2.011410e+00	1.000000	2.308161e+00	2.036689e+00	

(<https://cdn.analyticsvidhya.com/wp-content/uploa>


The numerical features are now centered on the mean (with a standard deviation of one) and the categorical features are one-hot encoded. (https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization)

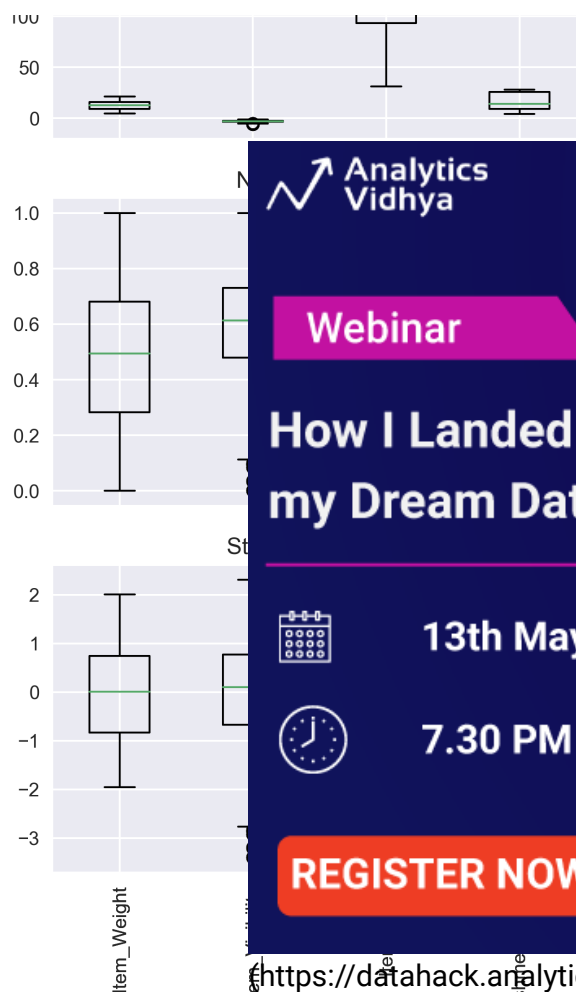
Comparing unscaled, normalized and standardized data

It is always great to visualize your data to understand the distribution present. We can see the comparison between our unscaled and scaled data using boxplots.

You can learn more about data visualization [here](https://www.analyticsvidhya.com/blog/tag/data-visualization/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization) (https://www.analyticsvidhya.com/blog/tag/data-visualization/?utm_source=blog&utm_medium=feature-scaling-machine-learning-normalization-standardization).

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).





Analytics Vidhya

Webinar

How I Landed my Dream Data Science Job?

13th May 2021

7.30 PM - 8.30 PM

REGISTER NOW

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

(https://cdn.analyticsvidhya.com/wp-content/uploads/2020/03/NormVsStand_box_plots-1.png)

You can notice how scaling the features brings everything into perspective. The features are now more comparable and will have a similar effect on the learning models.

Applying Scaling to Machine Learning Algorithms

It's now time to train some machine learning algorithms on our data to compare the effects of different scaling techniques on the performance of the algorithm. I want to see the effect of scaling on three algorithms in particular: K-Nearest Neighbours, Support Vector Regressor, and Decision Tree.

We use cookies on Analytics Vidhya websites to deliver our services, analyze web usage, and improve your experience on the site. By using Analytics Vidhya you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Accept

K-Nearest Neighbours

Like we saw before, KNN is a distance-based algorithm that is affected by the range of features. Let's see how it performs on our data, before and after scaling:

```

1  # training a KNN model
2  from sklearn.neighbors import KNeighborsRegressor
3  # measuring RMSE score
4  from sklearn.metrics import mean_squared_error
5
6  # knn
7  knn = KNeighborsRegressor(n_neighbors=7)
8
9  rmse = []
10
11 # raw, normalized and standardized training and testing data
12 trainX = [X_train, X_train_norm, X_train_std]
13 testX = [X_test, X_test_norm, X_test_std]
14
15 # model fitting and measuring RMSE
16 for i in range(len(trainX)):
17
18     # fit
19     knn.fit(trainX[i], y_train)
20     # predict
21     pred = knn.predict(testX[i])
22     # RMSE
23     rmse.append(np.sqrt(mean_squared_error(y_test, pred)))
24
25 # visualizing the result
26 df_knn = pd.DataFrame({'RMSE':rmse}, index=['Original', 'Normalized', 'Standardized'])
27 df_knn

```

view raw (<https://gist.github.com/aniruddha27/66119a2050fc808d2bdb7d4544ae75b6>
 /raw/f7d5d7854dfc734b910aefc9513ab7e3140c175e/NormalizationVsStandardization_4.py)
 NormalizationVsStandardization_4.py (https://gist.github.com/aniruddha27/66119a2050fc808d2bdb7d4544ae75b6#file-normalizationvsstandardization_4-py) hosted with ❤️ by GitHub (<https://github.com>)

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

	RMSE
Original	1319.283626
Normalized	1174.205859

Standardized 1183.448734

(https://cdn.analyticsvidhya.com/wp-content/uploads/2020/03/NormVsStand_knn.png).

You can see that scaling the features has brought down the RMSE. Normalized data performs a tad bit better than the standard data.

Note: I am measuring the RMSE here because this is a regression problem.

Support Vector Regressor

SVR (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-svm/>) is a distance-based algorithm. So let's check out whether

```

1  # training an SVR model
2  from sklearn.svm import SVR
3  # measuring RMSE score
4  from sklearn.metrics import mean_squared_error
5
6  # SVR
7  svr = SVR(kernel='rbf', C=5)
8
9  rmse = []
10
11 # raw, normalized and standardized training and testing data
12 trainX = [X_train, X_train_norm, X_train_stand]
13 testX = [X_test, X_test_norm, X_test_stand]
14
15 # model fitting and measuring RMSE
16 for i in range(len(trainX)):
17
18     # fit
19     svr.fit(trainX[i], y_train)
20
21     # predict on test data
22     pred = svr.predict(testX[i])
23     rmse.append(np.sqrt(mean_squared_error(y_test, pred)))
24
25 # RMSE
26 print('RMSE: ', rmse)
```



Webinar

How I Landed my Dream Data Science Job?



13th May 2021



7.30 PM - 8.30 PM

REGISTER NOW

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Accept

```

25 # visualizing the result
26 df_svr = pd.DataFrame({'RMSE':rmse},index=['Original','Normalized','Standardized'])
27 df_svr

```

view raw (https://gist.github.com/aniruddha27/a49f58527ef000949b28470a33a67eca94629e52f586cacab13d/NormalizaionVsStandarization_5.py) (https://gist.github.com/aniruddha27/a49f58527ef000949b28470a33a67eca94629e52f586cacab13d/NormalizaionVsStandarization_5.py) hosted with ❤ by GitHub

Original
Normalized
Standardized

(<https://cdn.analyticsvidhya.com/wp-content/uploads/2020/04/feature-scaling-machine-learn-1024x768px.jpg>)

We can see that scaling the features does bring down the RMSE. This means that the model performed better than the normalized data. Why do we need to scale the features?

The [sklearn documentation](https://scikit-learn.org/stable/modules/feature_scaling.html) (https://scikit-learn.org/stable/modules/feature_scaling.html) states that SVM

is centered around zero and variance is of the same order. This is because a feature with a variance greater than that of others prevents the estimator from learning from all the features. Great!

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

Decision Tree

We already know that a Decision tree is invariant to feature scaling. But I wanted to show a practical example of how it performs on the data:

```

1 # training a Decision Tree model
2 from sklearn.tree import DecisionTreeRegressor
3 # measuring RMSE score
4 from sklearn.metrics import mean_squared_error

```

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Accept

```

10
11 # raw, normalized and standardized training and testing data
12 trainX = [X_train,X_train_norm,X_train_stand]
13 testX = [X_test,X_test_norm,X_test_stand]
14
15 # model fitting and measuring RMSE
16 for i in range(len(trainX)):
17
18     # fit
19     dt.fit(trainX[i],y_train)
20     # predict
21     pred = dt.predict(testX[i])
22     # RMSE
23     rmse.append(np.sqrt(mean_squared_error(y_
24
25 # visualizing the result
26 df_dt = pd.DataFrame({'RMSE':rmse},index=['Or
27 df_dt

```

view raw ([https://gist.github.com/aniruddha27/6734a5390dc5b049c34d17ffa4889edda3b89c7827dcf97e5b498/Normal](https://gist.github.com/aniruddha27/6734a5390dc5b049c34d17ffa4889edda3b89c7827dcf97e5b498/Normal%20NormalizationVsStandarization_6.py)
NormalizationVsStandarization_6.py ([https://gist.github.com/aniruddha27/6734a5390dc5b049c34d17ffa4889edda3b89c7827dcf97e5b498/Normal](https://gist.github.com/aniruddha27/6734a5390dc5b049c34d17ffa4889edda3b89c7827dcf97e5b498/Normal%20NormalizationVsStandarization_6.py)
normalizationvsstandarization_6-py) hosted with ❤ by GitHub

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

RMSE

Original	1245.37439
Normalized	1245.37439
Standardized	1245.37439

(https://cdn.analyticsvidhya.com/wp-content/uploads/2020/03/NormVsStand_dt.png)

You can see that the RMSE score has not moved an inch on scaling the features. So rest assured when you are using tree-based algorithms on your data!


End Notes We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>). This tutorial covered the relevance of using feature scaling on your data and how normalization and standardization have varying effects on the working of machine learning algorithms.

As a next step, I encourage you to try out feature scaling, normalization or standardization? I recommend you to check out this article (<https://datahack.analyticsvidhya.com/contest/practice-problem-101/#feature-scaling-machine-learning>) with `utm_medium=feature-scaling-machine-learning-normalization` for continuity with this article. And don't forget to share

https://apps.apple.com/us/app/analytics-vidhya/id1444387141?referrer=utm_source=google&utm_medium=cpc&utm_campaign=MKT-Other-global-all-co-prtnr-py-PartBW&utm_term=analytics%20vidhya

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-first-job/?utm_source=blog&utm_medium=popup)

Scaling



We use cookies on Analytics

ya websites to deliver our services, a
/blog/2020/10/7-feature-engineering
e. By using Analytics Vidhya, you agr
techniques-machine-learning/

analyze web traffic, and improve your
/blog/2020/07/types-of-feature-
transformation-and-scaling/)

(<https://www.analyticsvidhya.com/blog/2020/12/feature-engineering-feature-improvements-scaling/>) We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

18 of 29

TAGS : [FEATURE SCALING \(HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/TAG/FEATURE-SCALING-MACHINE-LEARNING\)](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning/) [LEARNING \(HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/TAG/FEATURE-SCALING-MACHINE-LEARNING\)](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning/) [PYTHON \(HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/TAG/FEATURE-SCALING-MACHINE-LEARNING-PYTHON\)](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning/) [\(HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/TAG/LIVE-CODE\)](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning/) [\(HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/TAG/NORMALIZATION\)](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning/) [\(HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/TAG/NORMALIZATION\)](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning/) [\(HTTPS://WWW.ANALYTICSVIDHYA.COM/BLOG/TAG/STANDARDIZATION\)](https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning/)

Supervised Learning vs. Unsupervised Learning

(<https://www.analyticsvidhya.com/blog/2020/04/supervised-learning-vs-unsupervised-learning/>)

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

[PREVIOUS ARTICLE](#)

6 Open Source Data Science Projects to Make you Industry Ready!

(<https://www.analyticsvidhya.com/blog/2020/04/6-open-source-data-science-projects-industry-ready/>)



We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Aniruddha Bhandari (<https://www.analyticsvidhya.com/blog/author/aniruddha/>)

I am on a journey to becoming a data scientist. I love to unravel trends in data, visualize it and predict the future with ML algorithms! But the most satisfying part of this journey is sharing my learnings, from the challenges that I face, with the community to make the world a better place!

This article is quite old and you might not get a pr
this comment on Analytics Vidhya's [Discussion](#) p
queries resolved

16 COMMENTS



ALI

April 12, 2020 at 11:18 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-161122>)

Excelent article! Thank you very much for sharing. I
to fit the scaler on the training data and then use it t
you posted. Am I wrong? How would one "fit the sca
testing data"? Thanks a lot again

Analytics Vidhya

Webinar

How I Landed my Dream Data Science Job?

13th May 2021

7.30 PM - 8.30 PM

REGISTER NOW

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)



ANIRUDDHA BHANDARI

[Reply](#)

April 13, 2020 at 11:44 am (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-161126>)

Hi Ali

You fit the scaler on the training data so that it can calculate the necessary parameters, like mean and standard deviation for standardization, and store it for later use using the fit() method. Later you use the transform() function to apply the same transformation on both, train and test dataset.

I have used this approach for both, normalization and standardization, in the article in the gists "NormalizationVsStandarization_2.py" and "NormalizationVsStandarization_3.py" respectively. I hope this cleared your doubt.

Thank you for using cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our [Privacy Policy](#)

(<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).



SAHIL KAMBOJ

Accept

[Reply](#)

April 28, 2020 at 4:25 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-161126>)

[standardization/#comment-161208](#))

Good article! Thank you very much for sharing. I have one question.

What the difference between sklearn.preprocessing
sklearn.preprocessing.Normalizer?

When to use MinMaxScaler and when to Normalize?



ANIRUDDHA BHANDARI

April 29, 2020 at 2:15 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-161217>)

Hi

I hope MinMaxScaler is already clear from the article
Normalizer (<https://scikit-learn.org/stable/modules/sklearn.preprocessing.Normalizer.html#sklearn.preprocessing.Normalizer>)

The only difference is the way it computes the norm
row values i.e. each element of a row is normalized
elements in that row.

As mentioned in the documentation, it is useful in te
(<https://www.analyticsvidhya.com/blog/2020/02/quantile-normalization/>)

cosine similarity between the different sentences/documents in the dataset. Other than that, as I mentioned in
the article, there is no sure way to know which scaling technique should be used when. The best way is to create
multiple scaled copies of the data and then try them out and see which one gives the best result.

Hope this helps.



SUBHASH KUMAR NADAR

[Reply](#)

May 23, 2020 at 11:25 am (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-161548>)

Excellent article! Easy to understand and good coverage

One question: I see that there is a scale() funtion as well from sklearn and short description suggest it to be
similar to StandardScaler i.e. scaling to unit variance

I could not find more than this explanation. Please can you suggest which on to use which scenario?

Thanks in advance!

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your
experience on the site. By using Analytics Vidhya, you agree to our [Privacy Policy](#)

(<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).



ANIRUDDHA BHANDARI

[Reply](#)

May 24, 2020 at 12:41 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-161548>)

[standardization/#comment-161579](#))

Hi Subhash

I notice two differences between the two functions. your data along any axis. This means that you could wise, which is what happens in StandardScaler(). Se transform methods, so you cannot apply the same s I would suggest using the StandardScaler() function I hope this helps!



GOLLA KEDARKUMAR

May 24, 2020 at 9:56 am (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-161575>)

Hi ANIRUDDHA,

If we use the same scaler for train and testing, does need to use the mean of the data. If we take the mean the test data, right?



ANIRUDDHA BHANDARI

May 24, 2020 at 12:11 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-161578>)

Scaling your test data according to the train data makes sure that the test data is on the same scale as the training data on which our model was trained on. This way our model will be able to apply the learnings from the training dataset on the testing dataset, which is exactly what we want! If instead, we scale the test data differently, then our model might not be able to discern that difference, thereby giving us incorrect outputs. That way we will never know how well our model is performing. I hope this helps!



INAS

May 27, 2020 at 6:42 am (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-161630>)

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Accept

Analytics Vidhya

Webinar

How I Landed my Dream Data Science Job?

13th May 2021

7.30 PM - 8.30 PM

REGISTER NOW

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job?utm_source=blog&utm_medium=popup)

**SOUMADIP ROY**[Reply](#)

July 4, 2020 at 12:59 am (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-162136>)

This is an excellent write up. Thanks for this.

**HARSHVARDHAN BHATT**

July 5, 2020 at 3:03 am (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-162151>)

That graphs really helps in putting things in perspective.

**ARNOB**

July 5, 2020 at 4:13 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-162155>)

Hey bro! Great article. It covered a lots of topics that I needed. How can I check my data after normalization. You have a section on sklearn. But when I use it I get an error – "ValueError: Input data contains NaN"

Can you tell me how to check my data after normalization?

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

Thank you for your time.

**ANIRUDDHA BHANDARI**[Reply](#)

August 22, 2020 at 8:06 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-162806>)

Hi Arnob, glad you liked the article.

The command you are looking for is `df.describe()` not `pd.describe()`. Try using that, it should work.

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our [Privacy Policy](https://www.analyticsvidhya.com/privacy-policy/) (<https://www.analyticsvidhya.com/privacy-policy/>) and [Terms of Use](https://www.analyticsvidhya.com/terms/) (<https://www.analyticsvidhya.com/terms/>).

Accept

**DEEPS**[Reply](#)

July 15, 2020 at 7:59 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-162296>)

Excellent article !

**KUNAL**

July 31, 2020 at 1:19 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-162543>)

Thanks for Great Article..!!!

**ZINEB**

August 17, 2020 at 3:29 pm (<https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/#comment-162737>)

Thanks Bhandari.

Easy to understand and very helpful.

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Accept



Get your idea online.

in.godaddy.com

POPULAR POSTS

Understanding Delimiters in Pandas read_csv() Function (https://www.analyticsvidhya.com/blog/2021/04/delimiters-in-pandas-read_csv-function/)

40 Questions to test a Data Scientist on Clustering

(<https://www.analyticsvidhya.com/blog/2017/02/test-data-scientist-clustering/>)

How to Download Kaggle Datasets using Jupyter Notebook (<https://www.analyticsvidhya.com/blog/2021/04/how-to-download-kaggle-datasets-using-jupyter-notebook/>)

Python List Programs For Absolute Beginners (<https://www.analyticsvidhya.com/blog/2021/04/python-list-programs-for-absolute-beginners/>)

Commonly used Machine Learning Algorithms (with Python and R Codes) (<https://www.analyticsvidhya.com/blog/2017/09/common-machine-learning-algorithms/>)

30 Questions to test a data scientist on K-Nearest Neighbors (kNN) Algorithm

(<https://www.analyticsvidhya.com/blog/2017/09/30-questions-test-k-nearest-neighbors-algorithm/>)

45 Questions to test a data scientist on basics of Deep Learning (along with solution)

(<https://www.analyticsvidhya.com/blog/2017/01/must-know-questions-deep-learning/>)

Introductory guide on Linear Programming for (aspiring) data scientists (<https://www.analyticsvidhya.com/blog/2017/02/introductory-guide-on-linear-programming-explained-in-simple-english/>)

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy

(<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

Accept

Analytics Vidhya

Webinar

How I Landed my Dream Data Science Job?

13th May 2021

7.30 PM - 8.30 PM

REGISTER NOW



CAREER RESOURCES



16 Key Questions
Science (h
science-ca
utm_medium


NOVEMBER 20, 2020

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)



Here's What You Need to Know to Become a Data Scientist!
(https://www.analyticsvidhya.com/blog/2021/01/heres-what-you-need-to-know-to-become-a-data-scientist/?&utm_source=Blog&utm_medium=CareerResourceWidget)

JANUARY 22, 2021



These 7 Signs Show you have Data Scientist Potential!
(https://www.analyticsvidhya.com/blog/2020/12/these-7-signs-show-you-have-data-scientist-potential/?&utm_source=Blog&utm_medium=CareerResourceWidget)

DECEMBER 3, 2020



We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).
(https://www.analyticsvidhya.com/blog/2020/11/how-to-have-a-career-in-data-science-business-analytics/?&utm_source=Blog&utm_medium=CareerResourceWidget)

Webinar

How I Landed my Dream Data Science Job?

13th May 2021

7.30 PM - 8.30 PM

REGISTER NOW



NOVEMBER 26, 2020

Should I become a data scientist (or a business analyst)?

(<https://www.analyticsvidhya.com/blog/2020/11/become-data-scientist-b/>)
 utm_medium=blog

NOVEMBER 2

RECENT POSTS

Complete Guide on Encoding Numerical Features
[/blog/2021/05/complete-guide-on-encode-numerical-features/](https://www.analyticsvidhya.com/blog/2021/05/complete-guide-on-encode-numerical-features/)

MAY 12, 2021

How AI Supports Logistics Industry and Transportation
[/blog/2021/05/how-ai-supports-logistics-industry-transportation/](https://www.analyticsvidhya.com/blog/2021/05/how-ai-supports-logistics-industry-transportation/)

MAY 11, 2021

Making Programming with Date and Time, less painful
[/05/making-programming-with-date-and-time-less-painful/](https://www.analyticsvidhya.com/blog/2021/05/making-programming-with-date-and-time-less-painful/)

MAY 11, 2021

Is there any need of Deep Learning? (<https://www.analyticsvidhya.com/blog/2021/05/is-there-any-need-of-deep-learning/>)

MAY 11, 2021



Webinar

How I Landed my Dream Data Science Job?



13th May 2021




7.30 PM - 8.30 PM

REGISTER NOW




(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)



(https://bootcamp.analyticsvidhya.com/?utm_source=Blog&utm_medium=blog)

Accept

utm_medium=stickybanner1#av-roadmap)



Analytics Vidhya

Data Science Blogathon

- Assured INR 1,00,000 for every published article
- Total prizes worth INR 1 Lakh+
- A chance to win iPad(8th Gen)

(<https://www.analyticsvidhya.com/>)

8th April - 3rd May 21

Register Now

Download App (<https://play.google.com/store/apps/details?id=com.analyticsvidhya.android>)

Analytics Vidhya

About Us (<https://www.analyticsvidhya.com/about-me/>)

Our Team (<https://www.analyticsvidhya.com/about-me/team/>)

Careers (<https://www.analyticsvidhya.com/about-me/career-analytics-discussions/>)

Contact us (<https://www.analyticsvidhya.com/contact/>)

Visit us

in

(<https://www.linkedin.com/company/analyticsvidhya/>)

f

(<https://www.facebook.com/analyticsvidhya/>)

yt

(<https://www.youtube.com/channel/UCqANLNUzObA>)

tw

(<https://twitter.com/analyticsvidhya>)

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our Privacy Policy (<https://www.analyticsvidhya.com/privacy-policy/>) and Terms of Use (<https://www.analyticsvidhya.com/terms/>).

© Copyright 2013-2020 Analytics Vidhya



Analytics Vidhya

Webinar

How I Landed my Dream Data Science Job?

13th May 2021

7.30 PM - 8.30 PM

REGISTER NOW

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

Apply Jobs (<https://www.analyticsvidhya.com/jobs/>)

Companies

Post Jobs (<https://www.analyticsvidhya.com/corporate/>)

Trainings (<https://courses.analyticsvidhya.com/>)

Hiring Hackathons (<https://datahack.analyticsvidhya.com/>)

Advertising (<https://www.analyticsvidhya.com/contact/>)

☰

A dark blue rectangular banner for a webinar. In the top left corner is the Analytics Vidhya logo, which consists of a white line graph icon and the text 'Analytics Vidhya'. Below the logo, the word 'Webinar' is written in white inside a magenta parallelogram. The main title 'How I Landed my Dream Data Science Job?' is in large white font, underlined with a magenta line. Below the title, there are two rows of information: the first row has a calendar icon and the date '13th May 2021'; the second row has a clock icon and the time '7.30 PM - 8.30 PM'. At the bottom center is a red button with the text 'REGISTER NOW' in white. On the right side, there are faint, stylized illustrations of a laptop and a bar chart. In the bottom right corner, there is a small grid of white dots.

Analytics Vidhya

Webinar

**How I Landed
my Dream Data Science Job?**

 **13th May 2021**

 **7.30 PM - 8.30 PM**

REGISTER NOW

(https://datahack.analyticsvidhya.com/contest/how-i-landed-my-dream-data-science-job/?utm_source=blog&utm_medium=popup)

We use cookies on Analytics Vidhya websites to deliver our services, analyze web traffic, and improve your experience on the site. By using Analytics Vidhya, you agree to our [Privacy Policy](https://www.analyticsvidhya.com/privacy-policy/) (<https://www.analyticsvidhya.com/privacy-policy/>) and [Terms of Use](https://www.analyticsvidhya.com/terms/) (<https://www.analyticsvidhya.com/terms/>).

Accept